



ANSIBLE DYNAMIC INVENTORY

PREPARED BY :
SRINIVASA K



www.joinDevOps.com

Introduction

Ansible Dynamic Inventory is a powerful feature that allows you to generate your inventory of hosts on the fly, rather than maintaining a static inventory file. This approach enables you to integrate Ansible with various cloud providers, such as AWS or Azure, and manage your hosts more efficiently.

What is Dynamic Inventory?

Dynamic Inventory is a mechanism that allows Ansible to query an external source, such as a cloud provider's API, to gather information about the hosts you want to manage. This information is then used to create a temporary inventory of hosts that Ansible can use to run playbooks.

Benefits of Dynamic Inventory:

Dynamic Inventory provides several benefits, including:

- **Automated host management:** Dynamic Inventory allows you to automate the process of managing your hosts, reducing the need for manual intervention.
- **Integration with cloud providers:** Dynamic Inventory enables you to integrate Ansible with various cloud providers, making it easier to manage your hosts across different environments.
- **Improved scalability:** Dynamic Inventory allows you to manage a large number of hosts more efficiently, making it an ideal solution for large-scale environments.

Types of Ansible Inventory:

1. **Static Inventory:** Manually defined in an INI or YAML file. Suitable for small, stable environments.
2. **Dynamic Inventory:** Automatically fetched from external sources like cloud providers, databases, or APIs. Ideal for large-scale and cloud-based environments.

How Ansible Dynamic Inventory Works:

Ansible Dynamic Inventory is a feature that allows Ansible to fetch host data dynamically from external sources, eliminating the need for manual updates. Here's a step-by-step explanation of the process:

Step 1: External Source Connection

Ansible connects to an external source, such as:

- Cloud providers (AWS, Azure, GCP) Databases (MySQL, PostgreSQL)

Step 2: Plugin or Script Execution

Ansible executes a plugin or external script that generates the inventory. These plugins or scripts are specifically designed to fetch data from the connected external source.

Step 3: Data Fetching

The plugin or script fetches the required data from the external source. This data typically includes host information, such as IP addresses, hostnames, and group membership.

Step 4: Inventory Generation

Ansible generates the inventory based on the fetched data. The inventory is a JSON-formatted file that contains the host information.

Step 5: Real-time Updates

Ansible updates the inventory in real-time, eliminating the need for manual updates. This ensures that the inventory always reflects the current state of the external source.

Step 6: Playbook Execution

Ansible executes the playbook using the dynamically generated inventory. The playbook can then use the host information to perform various tasks, such as configuration management, deployment, and orchestration.

How to Use Ansible Dynamic Inventory:

1. Install Required Dependencies: For AWS

```
pip install boto3 botocore
```

2. Configure Dynamic Inventory for Cloud Providers

AWS Dynamic Inventory Configuration (aws_ec2.yml)

Create a file named `aws_ec2.yml` with the following content:

```
plugin: amazon.aws.aws_ec2
regions:
  - us-east-1
filters:
  instance-state-name: running
keyed_groups:
  - key: tags.Name
    prefix: tag
  - key: instance_type
    prefix: type
```

Run the inventory command to see the list of hosts:

```
ansible-inventory -i aws_ec2.yml --list
```

3. Running Playbooks with Dynamic Inventory

Use the inventory file with your playbooks. For example:

```
ansible-playbook -i aws_ec2.yml your_playbook.yml
```

Interview Questions & Answers

Q1: Your company has a multi-cloud environment. How would you manage inventory dynamically in Ansible?

A: I would configure dynamic inventory plugins for both AWS and Azure, ensuring Ansible fetches real-time host data from both cloud providers and integrates them into a single inventory.

Q2: How would you troubleshoot if Ansible Dynamic Inventory is not fetching the correct host details?

A: I would verify the external plugin configurations, test API connectivity, validate filter conditions, and use `ansible-inventory --list` to debug fetched data.

Q3: Your company frequently scales up and down its infrastructure. How would you ensure Ansible dynamically adapts to changes?

A: By leveraging Ansible Dynamic Inventory, Ansible can automatically detect new instances and remove terminated ones, ensuring an up-to-date infrastructure.

Q4: If an external API used for Dynamic Inventory is slow or unavailable, how can you mitigate playbook failures?

A: I would implement caching to store inventory data temporarily and set timeouts to avoid long delays when the API is unresponsive.

Q5: How can you group instances dynamically based on their tags in Ansible Dynamic Inventory?

A: I would use the `keyed_groups` option in the inventory configuration file to dynamically categorize hosts based on tags such as environment or instance type.

Q6: Can Ansible Dynamic Inventory be integrated into a CI/CD pipeline, and how?

A: Yes, it can be integrated by automating the inventory update before each deployment, ensuring playbooks always execute against the latest infrastructure state.

Conclusion:

Ansible Dynamic Inventory enhances infrastructure management by dynamically fetching host data from external sources like cloud providers, databases, and APIs. It automates inventory updates in real-time, eliminating manual intervention and improving efficiency. Ideal for multi-cloud or rapidly scaling environments, it simplifies management across platforms like AWS, Azure, and GCP. Dynamic Inventory allows for dynamic host grouping, real-time updates, and smooth handling of scaling challenges. Integrating it into CI/CD pipelines ensures playbooks always operate on the latest infrastructure, optimizing deployment and configuration management.